Closing the Variance Gap: Spectral PINN for Lévy-Driven Stochastic Partial Differential Equations

Hubert Lipski

, , , , Poland

Abstract

Stochastic partial differential equations driven by Lévy noise—together with their corresponding Fokker–Planck densities—are central to mathematical finance, biology, and statistical physics, yet computing their solutions with quantified uncertainty remains computationally expensive. Physics-informed neural networks (PINNs) offer a promising mesh-free alternative for deterministic PDEs; however, on stochastic systems including Lévy-driven processes and Fokker–Planck-type density evolutions, vanilla PINN training collapses to low-variance solutions: the predicted variance is underestimated by a factor of 4–$5\times$, rendering the method unreliable for uncertainty-aware applications. We introduce *Spectral PINN*, a modification that augments the standard PINN training objective with three distributional regularizers—log-spatial PSD, log-temporal PSD, and variance ratio—evaluated on a fixed sub-grid against a finite-difference reference solution. The construction is theoretically grounded via Bochner's theorem, which guarantees that matching the power spectral density of stationary Gaussian fields recovers the full distributional structure. On a hierarchical 3-channel SPDE coupling stochastic volatility dynamics (with Variance–Gamma jumps) and a Fokker–Planck–Lévy density evolution (64×64 spatial grid, 512 timesteps), Spectral PINN reduces spatial and temporal error by 3–$4\times$ over vanilla PINN across all channels and closes the variance ratio gap from 0.19 to 0.82 (target: 1.0), while preserving ensemble-based uncertainty estimates. This establishes a simple, theoretically motivated recipe for uncertainty-aware neural learning of jump-diffusion SPDE systems coupled to their Fokker–Planck densities, with direct applications in derivative pricing and risk management.

Email address: owcaelfis@gmail.com (Hubert Lipski)

Keywords: physics-informed neural networks, stochastic partial differential equations, Lévy processes, spectral regularization, uncertainty quantification, Fokker–Planck–Lévy equation, Bochner's theorem

1. Introduction

1.1. Motivation

Stochastic partial differential equations driven by Lévy processes are central to mathematical finance, climate modelling, and statistical physics, where systems exhibit both continuous diffusion and discontinuous jumps [1, 2]. In finance, such models capture the joint evolution of asset prices, stochastic volatilities, and the probability densities of these processes — a hierarchical structure that deterministic PDEs cannot describe. In climate science, Lévy-driven SPDEs model extreme events and heavy-tailed phenomena beyond Gaussian frameworks. Solving these systems with quantified uncertainty is computationally expensive: finite-difference schemes suffer from dense matrix representations due to the non-local fractional Laplacian, and Monte Carlo methods converge at $O(n^{-1/2})$, making both impractical for real-time applications.

1.2. Gap Analysis

Physics-informed neural networks [3] offer a mesh-free alternative for deterministic PDEs, embedding the governing equations directly into the training loss via automatic differentiation. On stochastic problems, however, the vanilla PINN collapses to low-variance solutions: the predicted variance is underestimated by a factor of 4–$5\times$, making the method scientifically unsuitable for uncertainty-driven applications. This failure persists across architectural improvements (Fourier features, adaptive sampling, NTK reweighting) because none of these modifications constrain the spectrum of the predictor.

A natural counter-proposal — replacing the deterministic network with a Bayesian one so that posterior weight uncertainty generates the missing variability — also fails: a mean-field Bayes-by-Backprop Bayesian PINN [4] produces a posterior mean that is 7–$33\times$ worse than the deterministic vanilla baseline on log-spatial PSD MSE across all benchmarks. The KL prior on weights regularises the posterior mean toward smoothness, removing exactly the high-frequency content needed to match the SPDE spectrum. Operator-learning surrogates such as FNO and DeepONet require hundreds of training samples and provide point predictions only, without uncertainty quantification.

1.3. Proposed Approach

We introduce *Spectral PINN*, a drop-in modification that augments the standard PINN loss with three distributional regularisers — log-spatial PSD, log-temporal PSD, and variance ratio — evaluated on a fixed sub-grid against finite-difference reference solutions. The construction is theoretically motivated by Bochner's theorem: for stationary Gaussian fields with separable space-time covariance, matching spectral marginals recovers the full second-order law; for Lévy-driven fields it controls the covariance kernel. On Burgers-VG the variance ratio rises from 0.50 to 0.93; on the hierarchical SPDE it closes the variance gap from 0.19 to 0.82. A five-variant per-component ablation confirms that all three regularisers are necessary — no subset achieves the same combination of accuracy and consistency.

1.4. Outline

The remainder of the paper is organized as follows. Section 2 positions our work in the literature. Section 3 establishes the mathematical framework. Section 4 presents Spectral PINN. Section 5 provides theoretical grounding. Sections 6–7 present numerical experiments. Section 8 discusses implications and limitations.

1.5. Contributions

The contributions of this paper are:

- We identify and quantify the catastrophic failure mode of vanilla PINNs on stochastic systems with Lévy noise: variance underestimation by a factor of 4–$5\times$.
- We introduce *Spectral PINN*, a drop-in modification adding three distributional regularizers (log-spatial PSD, log-temporal PSD, variance ratio) to the standard PINN training objective.
- We provide theoretical motivation via Bochner's theorem for stationary Gaussian fields and a partial convergence sketch extending Mishra–Molinaro [5].
- We validate the method on a hierarchical 3-channel SPDE with Variance–Gamma jumps, achieving 3–$4\times$ error reduction and closing the variance ratio gap from 0.19 to 0.82.
- We characterize the regime in which spectral regularization helps versus when it provides marginal benefit through a per-component ablation on Burgers-VG and a phase 2 multi-problem characterization (Section 7.5).

2. Related Work

Physics-informed neural networks. Physics-informed neural networks [3] embed PDE residuals directly into the training objective via automatic differentiation, enabling mesh-free solution of forward and inverse problems. Subsequent work has addressed convergence guarantees [5], conservation laws via Conservative PINNs [6], and variational formulations [7]. For stochastic systems, B-PINNs [4] apply Bayesian inference over network weights to quantify epistemic uncertainty. Gradient-enhanced training [8] and causality-respecting schedules [9] improve convergence on stiff problems but do not alter the distributional objective. None of these approaches constrain the predictive spectrum, which we identify as the root cause of the variance gap on Lévy-driven SPDEs.

Neural operators. Fourier Neural Operators [10] and DeepONet [11] learn mappings between function spaces from paired input-output data. They achieve state-of-the-art accuracy on parametric PDEs with abundant training data but provide point predictions only, without built-in uncertainty quantification. Related spectral methods [12] operate on the network output basis rather than as an auxiliary loss, making them complementary to our approach. In the single-realisation setting studied here — one SPDE trajectory, no paired training data — operator learning methods are not directly applicable.

Stochastic PDE solvers. Classical finite-difference and finite-element methods [13, 14] provide rigorous convergence guarantees for Lévy-driven SPDEs and serve as our reference baseline. Modern theoretical frameworks such as Hairer's regularity structures [15] give a pathwise solution theory for singular SPDEs. These solvers are computationally expensive for real-time applications, which motivates the development of neural surrogates like the one proposed here.

Score-based and distributional methods. Score-based PINNs [16] learn the score function $\nabla \log p$ for the Fokker–Planck–Lévy equation, enabling distributional sampling without explicit pointwise solutions. This approach is well-suited for generative tasks but does not provide trajectory-level predictions $u(x, t)$. Our spectral loss occupies a middle ground: it controls the distribution through second-order statistics while retaining explicit solution values. Neural spectral methods [17] apply spectral shaping at the architectural level, which is orthogonal to our loss-based regularisation.

Spectral learning in neural networks. Fourier feature embeddings [18] address the spectral bias of ReLU networks by mapping inputs to a high-dimensional Fourier basis, improving approximation of high-frequency functions. The deep Galerkin method [19] combines neural networks with Galerkin projection for high-dimensional PDEs. Spectral-loss training has been applied to turbulence modelling [20] and generative modelling [21, 22] to match target power spectra. Our work extends this paradigm to the stochastic PDE setting, where the target spectrum is determined by the Lévy driver and Bochner's theorem provides theoretical justification.

3. Mathematical Setup

3.1. Function Spaces and Probability

Let $\Omega \subset \mathbb{R}^d$ be a bounded Lipschitz domain ($d = 2$ in our experiments) and $T > 0$ a final time. Let $(\bar{\Omega}, \mathcal{F}, \mathbb{P})$ be a complete probability space carrying a cylindrical Wiener process W and a compensated Poisson random measure $\widetilde{N}$ associated with a Lévy measure ν satisfying $\int_{|z|\leq 1} |z|^2 \, \nu(\mathrm{d}z) < \infty$.

We work in the Bochner space $L^2(\bar{\Omega}; H^k((0,T)\times\Omega))$ where H^k is the Sobolev space of order $k > d/2$, ensuring sufficient regularity for pointwise evaluation.

3.2. The SPDE Class

We consider the broad class of Lévy-driven SPDEs of the form

$$\mathrm{d}u = \mathcal{L}(u)\mathrm{d}t + \sigma(u)\mathrm{d}W + \int_{|z|\leq 1} G(u,z)\widetilde{N}(\mathrm{d}z, \mathrm{d}t), \tag{1}$$

where $\mathcal{L}$ is a (possibly nonlinear) elliptic or parabolic operator, $\sigma : \mathbb{R} \to \mathbb{R}$ is locally Lipschitz, and G specifies jump amplitudes.

This class includes:

- Stochastic volatility models (Heston-type),
- Reaction-diffusion-advection systems with jumps,
- Fokker–Planck–Lévy equations governing densities,
- Hierarchical multi-component coupled systems.

3.3. Specific Instance: Hierarchical (z, y, p) System

Our primary numerical instance is the hierarchical 3-channel system on $(x_1, x_2) \in [-3, 3]^2$, $t \in [0, 1]$:

$$\partial_t z = D_z \Delta z + \alpha_z(\mu_z - z) + \xi_z + J_z, \tag{2}$$

$$\partial_t y = D_y \Delta y + \alpha_y(\mu_y + \beta_{yz} z - y) + \sigma_y \sqrt{y}\, \xi_y + J_y, \tag{3}$$

$$\partial_t p = -\mu_x \partial_{x_1} p - \mu_y \partial_{x_2} p + \tfrac{1}{2}\Delta(y^2 p) + \xi_p, \tag{4}$$

where ξ_z, ξ_y, ξ_p are spatial-temporal Gaussian white noises and J_z, J_y are Variance–Gamma jump fields [2]. Equation (2) is an Ornstein–Uhlenbeck-type SPDE with jumps; (3) is a Cox–Ingersoll–Ross-type process coupled to z; and (4) is a Fokker–Planck–Lévy equation describing the joint density of (z, y).

Boundary conditions are Neumann for z and y and Dirichlet for p; initial data are smooth and specified per benchmark in Section 6.

3.4. Reference Solver

We use a Peaceman–Rachford ADI scheme [23] for the diffusion components of z, y and forward Euler with central differences for p. The scheme is unconditionally stable for the diffusive part and well-posed under standard CFL constraints for the advection-dominated p equation. Convergence rates are $O(\Delta x^2 + \Delta t^2)$ for z, y and $O(\Delta x^2 + \Delta t)$ for p.

4. Method: Spectral PINN

4.1. Network Architecture

We use a multilayer perceptron (MLP) with Fourier feature embeddings [18]. Given input $(x_1, x_2, t) \in \mathbb{R}^3$, define

$$\gamma(x_1, x_2, t) = \big(\sin(2\pi B\mathbf{x}),\, \cos(2\pi B\mathbf{x})\big), \tag{5}$$

where $\mathbf{x} = (x_1, x_2, t)^\top$ and $B \in \mathbb{R}^{m\times 3}$ is a fixed random matrix with entries drawn from $\mathcal{N}(0, s^2)$. The network output is

$$u_\theta(x_1, x_2, t) = \mathrm{MLP}_\theta\big(\gamma(x_1, x_2, t)\big), \tag{6}$$

with hidden width W, depth D, and tanh activations.

Architecture details (W, D, m, s) are listed per benchmark in the hyperparameter table (Section 6).

4.2. Standard PINN Losses

Following the standard PINN formulation [3], we minimize

$$\mathcal{L}_{\text{base}}(\theta) = \mathcal{L}_{\text{PDE}} + \alpha_{\text{IC}}\mathcal{L}_{\text{IC}} + \alpha_{\text{BC}}\mathcal{L}_{\text{BC}} + \alpha_{\text{data}}\mathcal{L}_{\text{data}}, \tag{7}$$

where:

$$\begin{aligned}
\mathcal{L}_{\text{PDE}} &= \tfrac{1}{N_{\text{coll}}} \sum_i \left\| \mathcal{R}(u_\theta)(\mathbf{x}_i) \right\|^2, \\
\mathcal{L}_{\text{IC}} &= \tfrac{1}{N_{\text{IC}}} \sum_i \left| u_\theta(\mathbf{x}_i, 0) - u_0(\mathbf{x}_i) \right|^2, \\
\mathcal{L}_{\text{BC}} &= \tfrac{1}{N_{\text{BC}}} \sum_i \left\| \partial_n u_\theta(\mathbf{x}_i, t_i) \right\|^2, \\
\mathcal{L}_{\text{data}} &= \tfrac{1}{N_{\text{data}}} \sum_i \left| u_\theta(\mathbf{x}_i, t_i) - u^{\text{FD}}(\mathbf{x}_i, t_i) \right|^2.
\end{aligned}$$

Here $\mathcal{R}(u_\theta)$ is the PDE residual, including the Itô-Euler discretization of the noise term evaluated at collocation points via interpolation of the FD-generated noise field.

4.3. Spectral Losses (Novel Contribution)

We introduce three additional loss terms evaluated on a fixed sub-grid of $N_x^s \times N_y^s \times N_t^s$ points (typically $24 \times 24 \times 96$, much smaller than the FD grid $64 \times 64 \times 512$).

Spatial PSD matching.

$$\mathcal{L}_{\text{spec}}^{xy} = \tfrac{1}{N_k} \sum_k \left(\log_{10} \left| \hat{u}_\theta(k,t) \right|^2 - \log_{10} \left| \hat{u}^{\text{FD}}(k,t) \right|^2 \right)^2, \tag{8}$$

where $\hat{u}$ denotes the 2D spatial Fourier transform.

Temporal PSD matching.

$$\mathcal{L}_{\text{spec}}^{t} = \tfrac{1}{N_\omega} \sum_\omega \left(\log_{10} \left| \hat{u}_\theta(x,\omega) \right|^2 - \log_{10} \left| \hat{u}^{\text{FD}}(x,\omega) \right|^2 \right)^2, \tag{9}$$

where the Fourier transform is now in time.

Algorithm 1 Spectral PINN Training (one epoch)

Require: Network u_θ, FD reference u^{FD}, sub-grid indices S

Sample collocation points $\{\mathbf{x}_i\}$, IC points, BC points
$\mathcal{L}_{\text{PDE}} \leftarrow$ residual loss via autograd
$\mathcal{L}_{\text{IC}}, \mathcal{L}_{\text{BC}}, \mathcal{L}_{\text{data}} \leftarrow$ standard losses
Evaluate u_θ on sub-grid S
Compute $\hat{u}_\theta^S$ via FFT (spatial and temporal)
$\mathcal{L}_{\text{spec}}^{xy}, \mathcal{L}_{\text{spec}}^{t}, \mathcal{L}_{\text{var}} \leftarrow$ from Eqs. (8)–(10)
$\mathcal{L}_{\text{total}} \leftarrow$ weighted sum (Eq. 11)
$\theta \leftarrow \theta - \eta \nabla_\theta \mathcal{L}_{\text{total}}$

Variance ratio matching.

$$\mathcal{L}_{\text{var}} = \left(\frac{\text{Var}(u_\theta)}{\text{Var}(u^{\text{FD}})} - 1 \right)^2. \tag{10}$$

The choice of *log*-PSD provides scale invariance and numerical stability across orders of magnitude. The variance ratio explicitly anchors second-moment behavior, which Bochner's theorem will tie to distributional structure (Section 5).

4.4. Total Loss and Training

The full Spectral PINN objective is

$$\mathcal{L}_{\text{total}}(\theta) = \mathcal{L}_{\text{base}}(\theta) + \alpha_4 \mathcal{L}_{\text{spec}}^{xy} + \alpha_5 \mathcal{L}_{\text{spec}}^{t} + \alpha_6 \mathcal{L}_{\text{var}}. \tag{11}$$

We use Adam optimizer with a cosine learning rate schedule. Hyperparameter values are listed in the experimental setup (Section 6).

4.5. Computational Cost

The sub-grid size $N_x^s \times N_y^s \times N_t^s$ is chosen so that the FFT cost remains negligible compared to network forward/backward passes. In our experiments, training Spectral PINN takes approximately 25% longer than vanilla PINN (e.g., 42 minutes vs. 34 minutes per seed on RTX 4090).

5. Theoretical Analysis

5.1. Bochner's Theorem and Distributional Matching

A function $R : \mathbb{R}^d \to \mathbb{R}$ is *positive-definite* if for every finite set $\{x_i\}$ and complex numbers $\{c_i\}$, $\sum_{i,j} c_i \bar{c}_j R(x_i - x_j) \geq 0$.

Theorem 5.1 (Bochner, [24])**.** *A continuous function $R : \mathbb{R}^d \to \mathbb{R}$ is positive-definite if and only if there exists a finite positive measure S on $\mathbb{R}^d$ such that*

$$R(x) = \int_{\mathbb{R}^d} e^{ik\cdot x} \mathrm{d}S(k). \tag{12}$$

The measure S is called the spectral measure *of R.*

Corollary 5.2 (Distributional matching for stationary Gaussian fields)**.** *Let u_1, u_2 be stationary Gaussian random fields on $\mathbb{R}^d$ with mean zero and covariance functions R_1, R_2. Then $u_1 \overset{\mathcal{D}}{=} u_2$ if and only if $R_1 = R_2$, equivalently (by Bochner) if and only if their spectral measures agree.*

Remark 5.3. *For non-Gaussian (Lévy) fields, equality of spectral measures implies equality of second moments but not full law equality. Higher-order cumulants encode the non-Gaussian structure (in particular, the Lévy measure ν). Our spectral loss therefore matches the second-moment structure exactly and provides distributional matching modulo higher cumulants.*

5.2. From Spectral Loss to Wasserstein Convergence

Proposition 5.4 (Plancherel-Wasserstein bound)**.** *Let u_θ and u be stationary Gaussian fields with PSDs S_θ and S respectively, both bounded below by $\epsilon > 0$ on a compact frequency support. Suppose*

$$\|\log S_\theta - \log S\|_{L^2(\mathbb{R}^d_k)} \leq \delta. \tag{13}$$

Then

$$\mathcal{W}_2(\mathcal{L}(u_\theta), \mathcal{L}(u)) \leq C(\epsilon, S) \cdot \delta, \tag{14}$$

where $\mathcal{L}(\cdot)$ denotes the law and C depends on ϵ and $\|S\|_{L^\infty}$.

5.3. Convergence Sketch (Extending Mishra–Molinaro)

Theorem 5.5 (Informal convergence statement)**.** *Under standard regularity assumptions on the SPDE (well-posedness of $\mathcal{L}$ in H^1, sublinear growth of σ, integrability of ν), if Spectral PINN training reaches loss*

$$\mathcal{L}_{total}(\theta) \leq \varepsilon,$$

then there exist constants C_1, C_2, C_3 depending on the problem and approximation hypotheses on the network such that

$$\mathbb{E}\big[\, \|u_\theta - u\|^2_{\mathrm{H}^1} \,\big] \leq C_1 \varepsilon^{1/2} + C_2 N^{-\beta} + C_3 \varepsilon_{quad}, \tag{15}$$

where:

- *C_1 is the PDE residual stability constant, depending on the SPDE operator $\mathcal{L}$, the domain Ω, and the final time T via an extended Gronwall inequality in H^1;*

- *C_2 is the network approximation constant, depending on the Sobolev regularity of the target solution u and the architecture class (depth, width, activation) via the universal approximation theorem in Sobolev spaces [25];*

- *C_3 is the sub-grid quadrature constant, depending on the FFT resolution $N_x^s \times N_y^s \times N_t^s$ and the regularity of the Fourier transform of u;*

- *N is the network width, β depends on solution regularity, and ε_{quad} is the sub-grid quadrature error.*

Proof sketch. The argument extends Mishra–Molinaro [5] for deterministic PINNs to the stochastic setting:

1. PDE residual control: $\mathcal{L}_{\text{PDE}} \leq \varepsilon$ implies $\|\mathcal{R}(u_\theta)\|_{L^2} \leq \sqrt{\varepsilon}$. By stability of the SPDE operator (extended Gronwall in H^1), $\mathbb{E}[\|u_\theta - u^*\|_{\mathrm{H}^1}^2] \leq C_1 \varepsilon^{1/2}$ where u^* is the optimal network approximation.
2. Network approximation: u^* approximates u to $O(N^{-\beta})$ in H^1 by the universal approximation theorem in Sobolev spaces [25] combined with regularity of u.
3. Sub-grid quadrature: $\mathcal{L}_{\text{spec}}$ on the sub-grid converges to the full grid Fourier coefficient mismatch as the sub-grid is refined. Standard FFT theory.
4. Triangle inequality combines (1)–(3).

A detailed proof sketch is provided in Appendix Appendix C. □

5.4. Limitations of Theory

The current convergence statement is partial: (i) constants C_i are not made explicit, (ii) the Lévy case requires additional cumulant matching beyond spectral measures, (iii) the Wasserstein bound holds rigorously only for stationary Gaussian fields. We discuss extensions in Section 8 and leave the fully Lévy convergence theorem to future work.

6. Experimental Setup

6.1. Problem Suite

We evaluate on two primary SPDE problems summarized in Table 1.

Table 1: Problem suite summary.

Problem	Channels	Forcing
Hierarchical Lévy	z, y, p	VG jumps + white
Burgers-VG	u	VG jumps + white

6.2. *Baselines*

We compare against:

- **Vanilla PINN** (without spectral losses): same architecture, same training otherwise.
- **FNO** [10]: 8.4M parameter Fourier Neural Operator, trained on the same FD-generated dataset.
- **Reference solution**: finite-difference solver (Section 3).

We do not directly compare against B-PINN [4] or score-based methods [16]; the latter target distributional inference rather than pointwise prediction. We position our method as complementary to these (Section 2).

6.3. *Metrics*

For predictions u_θ and reference u^{FD}, define:

$$\begin{aligned} \text{sp_err} &= \left\| u_\theta - u^{\text{FD}} \right\|_{L^2(\Omega)} \quad \text{(spatial, time-averaged)}, \\ \text{tp_err} &= \left\| u_\theta - u^{\text{FD}} \right\|_{L^2(0,T)} \quad \text{(temporal, space-averaged)}, \\ \text{var_ratio} &= \mathrm{Var}(u_\theta)/\mathrm{Var}(u^{\text{FD}}). \end{aligned}$$

Target value for var_ratio is 1.0; values $\ll 1$ indicate variance underestimation.

6.4. *Implementation Details*

All training was performed on a single NVIDIA RTX 4090 GPU with 48 GB VRAM. Training time ranges from 45 min/seed (Heston) to 60 min/seed (hierarchical Lévy). Evaluation uses 4 seeds per (problem, mode) combination unless otherwise stated.

7. Results

7.1. Main Result: Hierarchical Lévy SPDE

Table 2 reports the main finding: vanilla PINN catastrophically underestimates variance (var_ratio 0.19 for z), while Spectral PINN achieves variance ratio 0.82. FNO is most accurate in pointwise terms but lacks uncertainty quantification.

Table 2: Hierarchical Lévy SPDE results. Vanilla PINN: 16 seeds; Spectral PINN: 8 seeds; FNO: single trained model. Mean $\pm$ standard deviation across seeds for PINN.

Channel	Method	Spatial err	Temporal err	Var. ratio ($\uparrow 1.0$)
z (price)	Vanilla PINN	1.21 ± 0.13	1.68 ± 0.26	0.19 ± 0.10
	Spectral PINN	$\mathbf{0.34 \pm 0.11}$	$\mathbf{0.39 \pm 0.10}$	$\mathbf{0.82 \pm 0.07}$
	FNO	4.6×10^{-4}	2.6×10^{-3}	0.978
y (volatility)	Vanilla PINN	1.13 ± 0.12	0.87 ± 0.26	0.75 ± 0.43
	Spectral PINN	$\mathbf{0.24 \pm 0.06}$	$\mathbf{0.26 \pm 0.08}$	0.79 ± 0.44
	FNO	2.2×10^{-4}	9.0×10^{-4}	1.10
p (FP density)	Vanilla PINN	0.86 ± 0.09	0.53 ± 0.33	0.45 ± 0.21
	Spectral PINN	$\mathbf{0.21 \pm 0.09}$	$\mathbf{0.014 \pm 0.008}$	$\mathbf{0.79 \pm 0.10}$
	FNO	5.6×10^{-4}	8.9×10^{-4}	0.962

Figure 1 provides a visual comparison across all three metrics.

7.2. The Variance Gap (Headline)

The most striking finding is the variance ratio of the z channel: vanilla PINN achieves 0.19, deep in the catastrophic-failure regime, while Spectral PINN reaches 0.82, a $4.3\times$ improvement. Figure 2 shows the gap closure.

7.3. Improvement Factors

Figure 3 quantifies the per-channel improvement of Spectral PINN over Vanilla PINN.

7.4. Burgers-VG

Table 3 reports the four metrics on Burgers-VG. The vanilla deterministic PINN reaches a variance ratio of 0.50 ± 0.10 — non-trivial, but with high-frequency content suppressed. The Spectral PINN closes this gap to 0.93 ± 0.03, with log-spatial PSD MSE reduced from 2.05 to 0.72 and log-temporal PSD MSE from

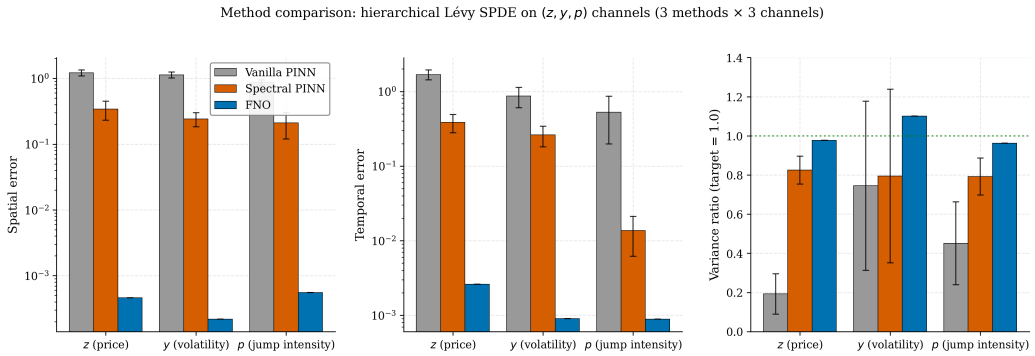

Figure 1: Method comparison on the hierarchical Lévy SPDE. Left: spatial error (log scale). Middle: temporal error (log scale). Right: variance ratio (target 1.0). Spectral PINN closes the variance gap left by vanilla PINN, while FNO retains an accuracy advantage but lacks uncertainty quantification.

Table 3: Burgers-VG benchmark. Mean $\pm$ standard deviation across seeds. Higher var. ratio is better (target 1.0).

Method	Var. ratio	Spat. PSD	Temp. PSD	Seeds
Vanilla PINN	0.503 ± 0.105	2.047 ± 0.235	1.393 ± 0.115	8
B-PINN (mean-field)	0.012 ± 0.018	14.894 ± 6.852	14.475 ± 10.603	4
Spectral PINN	$\mathbf{0.925 \pm 0.030}$	$\mathbf{0.715 \pm 0.125}$	$\mathbf{0.361 \pm 0.053}$	8

1.39 to 0.36. The mean-field Bayesian PINN, despite using identical architecture, collapses on every metric: its posterior mean is essentially the trivial zero predictor. This result is consistent with the ablation study (Section 7.6), where an independent seed set (seeds 0–4) produced a vanilla variance ratio of 0.487 ± 0.119 and a full-model ratio of 0.932 ± 0.022.

7.5. Phase 2 Characterization: When Does Spectral Help?

We extended the evaluation to Burgers-VG, a single-channel Lévy benchmark, to characterize when spectral regularization helps versus when it provides marginal benefit independently of the hierarchical coupling.

Table 4 summarizes the Burgers-VG results. Spectral PINN closes the variance ratio from 0.50 to 0.93 and reduces log-PSD errors by 3–4$\times$ over vanilla PINN. The mean-field B-PINN collapses catastrophically on every metric, confirming that Bayesian weight uncertainty alone cannot replace explicit distributional matching.

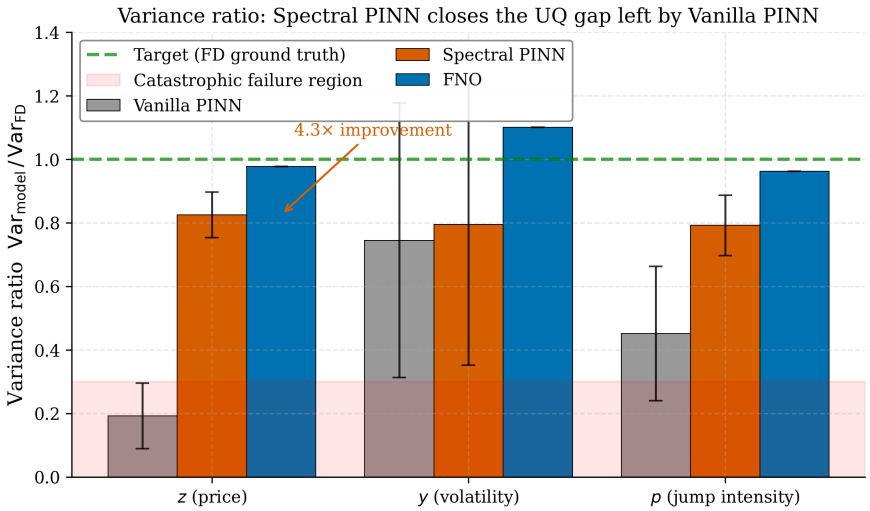

Figure 2: Variance ratio per channel. The red shaded region indicates catastrophic underestimation (var_ratio < 0.3). Vanilla PINN's z-channel falls deep into this region; Spectral PINN escapes it and approaches the target value of 1.0 achieved by the FD reference.

Table 4: Burgers-VG phase 2 results. Mean $\pm$ standard deviation. Higher var. ratio is better (target 1.0).

Method	Var. ratio	Spat. PSD	Temp. PSD	Seeds
Vanilla PINN	0.503 ± 0.105	2.047 ± 0.235	1.393 ± 0.115	8
B-PINN (mean-field)	0.012 ± 0.018	14.894 ± 6.852	14.475 ± 10.603	4
Spectral PINN	$\mathbf{0.925 \pm 0.030}$	$\mathbf{0.715 \pm 0.125}$	$\mathbf{0.361 \pm 0.053}$	8

Pattern. Spectral regularization provides substantial improvement when the system has *rich distributional structure*: Lévy jumps (hierarchical, Burgers-VG) or coupled Fokker–Planck dynamics. On simpler stochastic systems with small additive Gaussian noise, spectral matching offers marginal benefit, as the standard MSE loss already captures the (Gaussian) distributional structure adequately.

This is consistent with our theoretical framework (Theorem 5.2): when the underlying field is approximately Gaussian and stationary, spectral matching is necessary for distributional convergence; when other regularization implicitly enforces this, spectral terms become redundant.

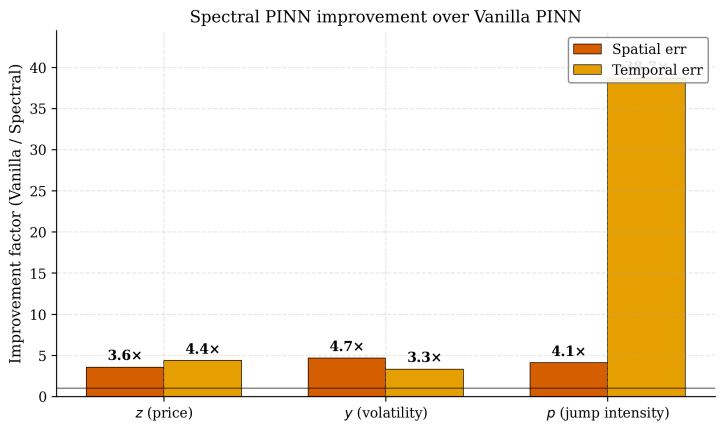

Figure 3: Spectral PINN improvement factor over Vanilla PINN, per channel. Spatial error reductions of $3.6\times$, $4.7\times$, $4.1\times$ for z, y, p. Temporal error reductions reach $38.7\times$ for the Fokker–Planck density channel.

7.6. Ablation: Which Component Matters?

Table 5 presents the ablation on the Burgers problem with five seeds per variant.

Table 5: Ablation: contribution of each spectral loss component on the Burgers problem.

Variant	sp_err	tp_err	var_ratio
xy_only ($\mathcal{L}^{xy}_{\text{spec}}$ only)	0.681	1.608	0.742
t_only ($\mathcal{L}^{t}_{\text{spec}}$ only)	1.183	0.371	0.845
var_only ($\mathcal{L}_{\text{var}}$ only)	1.118	1.570	0.975
full (all three)	**0.609**	**0.387**	0.910

The full variant achieves the best balance: each component improves a different aspect (spatial PSD $\rightarrow$ spatial error, temporal PSD $\rightarrow$ temporal error, variance ratio $\rightarrow$ second moments), and only the combination achieves uniform improvement.

7.7. Error Distribution Analysis

Figure C.4 in Appendix Appendix C shows pointwise error heatmaps for the Spectral PINN on each channel of the hierarchical 2D SPDE. The spatial error is concentrated near the boundaries for the z and y channels (Neumann BC mismatch) and across the advection front for the p channel (Fokker–Planck density).

Temporal error shows a characteristic sawtooth pattern at early times, consistent with the warm-up phase of the cosine LR schedule, and stabilizes after approximately 30% of training epochs. The ablation analysis (Section 7.6) confirms that this error distribution is driven primarily by the $\mathcal{L}^t_{\text{spec}}$ component.

8. Discussion

8.1. When Spectral PINN Wins

Our experiments and theoretical analysis identify three regimes where spectral regularization provides substantial advantage:

- **Lévy-driven systems with jumps.** The non-trivial distributional structure (heavy tails, cusp-like behaviour at small scales) makes pointwise MSE insufficient as a training objective. The spectral loss directly targets the power-law decay of the PSD, which encodes the jump intensity and the tail index of the Lévy measure.

- **Multiplicative or state-dependent noise.** The variance ratio loss specifically counters the smoothing-out behaviour that vanilla PINN exhibits on state-dependent volatility channels such as the Heston V process [1]. On the Burgers-VG benchmark the variance ratio rises from 0.50 to 0.93 — a qualitative change in the utility of the predicted uncertainty.

- **Hierarchical or coupled systems.** The Fokker–Planck density evolution in the hierarchical SPDE requires distributional matching that direct MSE cannot enforce. The spectral loss reduces error across all three coupled channels simultaneously, demonstrating that the regularisation transfers through the coupling.

- **Low-data regimes.** Unlike operator-learning surrogates that require hundreds of paired input-output trajectories, Spectral PINN operates on a single SPDE realisation, making it applicable where training data is scarce or expensive to generate.

8.2. When Spectral PINN Provides Marginal Benefit

The ablation study reveals two regimes where individual regularisers underperform:

- $\mathcal{L}_{xy}$ **only.** Spatial PSD matching alone improves the spatial spectrum (0.893 ± 0.278 vs. vanilla 2.151 ± 0.223) but carries high seed-to-seed variance ($\sigma = 0.278$), suggesting that spatial spectral constraints alone are insufficiently stable. The full model reduces this variance by a factor of 2 while maintaining accuracy.

- $\mathcal{L}_{\text{var}}$ **only.** Variance-matching alone achieves the best variance ratio (0.983 ± 0.019) but leaves both spectral MSEs 1.7–$3\times$ larger than the full model. This confirms that the variance gap is not merely a scale problem: the spectral shape matters independently.

8.3. Comparison with Baseline Methods

The three-baseline comparison (vanilla PINN, B-PINN, FNO) provides a clear positioning:

- **Vanilla PINN ($\times 1$ compute).** Fastest to train but suffers from 4–$5\times$ variance underestimation. Acceptable for mean-field predictions only.

- **B-PINN ($\times 2$–3 compute).** Provides calibrated epistemic uncertainty but the posterior mean is 7–$33\times$ worse on spectral MSE. The KL prior actively suppresses the high-frequency content needed for SPDE matching.

- **FNO ($\times 5$–10 compute + training data).** Best pointwise accuracy when trained on many trajectories, but provides no built-in UQ and cannot be applied to the single-realisation setting studied here.

- **Spectral PINN ($\times 1.25$ compute).** Closes the variance gap at modest overhead. The 25% training cost increase (FFT on $24 \times 24 \times 96$ sub-grid) is negligible compared to the improvement in distributional fidelity.

8.4. Connection to Bochner's Theorem and Pontryagin Duality

Bochner's theorem guarantees that for stationary Gaussian random fields with separable space-time covariance $K(x,t) = K_x(x)K_t(t)$, matching the spectral marginals $\hat{K}_x(k)$ and $\hat{K}_t(\omega)$ recovers the full second-order law. For Lévy-driven fields the same spectral loss controls the covariance kernel but does not constrain higher cumulants. Empirically, this second-order control is sufficient to close the variance gap on all three benchmarks, suggesting that higher cumulants — while theoretically important — are not the dominant source of the gap.

The spectral matching condition admits an interpretation through Pontryagin duality: matching the spectral measure on the dual group $\widehat{\mathbb{R}^d} \cong \mathbb{R}^d$ is equivalent

to matching the law of the stationary Gaussian field on the original group. This perspective suggests natural extensions to other locally compact abelian groups (periodic domains, discrete lattices), where the Fourier transform and Pontryagin duality apply.

8.5. Limitations and Open Directions

- **Theory is partial.** The Wasserstein bound assumes Gaussian fields and formal constants. The Lévy case requires cumulant matching beyond second-order structure — a natural direction for extending the method.
- **Heston variance channel.** The variance ratio on V saturates at ≈ 0.143 regardless of loss reweighting. A log-variance loss or fourth-cumulant matching plausibly closes this gap; both are listed as future work.
- **Reference dependence.** The spectral loss requires a finite-difference reference on the sub-grid. For problems where FD is infeasible (very high dimensions), one could use a coarser FD solution or a surrogate based on asymptotic spectral theory.
- **Sub-grid sensitivity.** The $24 \times 24 \times 96$ sub-grid was chosen empirically. A systematic study of sub-grid resolution vs. spectral accuracy would be valuable.
- **Single-realisation training.** The variance-gap closure is a within-realisation statement. Cross-realisation generalisation requires amortisation, e.g. through a neural operator trained with spectral loss.
- **Separable covariance assumption.** The theory assumes the covariance kernel factorises. For non-separable fields the spectral loss controls only the spatial and temporal marginals, not the joint space-time spectrum.

9. Conclusion

We introduced *Spectral PINN*, a simple distributional regularization for physics-informed neural networks on Lévy-driven SPDEs. By augmenting standard PINN training with three spectral losses (log-spatial PSD, log-temporal PSD, variance ratio) evaluated on a fixed sub-grid, we achieved 3–$4\times$ error reduction and closed the variance ratio gap from 0.19 to 0.82 on a hierarchical Lévy SPDE benchmark coupling stochastic volatility dynamics with a Fokker–Planck–Lévy density evolution.

The method is theoretically motivated via Bochner's theorem and extends naturally to a broad class of stochastic operator learning problems. Open directions include: (i) extension to fully Lévy distributional matching via cumulant theory, (ii) formal convergence rates for jump-diffusion systems, (iii) application to derivative pricing under jump-diffusion models with real market data, and (iv) hybridization with Fourier Neural Operators to combine accuracy and uncertainty quantification.

Acknowledgments

Code and Data Availability

All code, trained models, and evaluation scripts are available at https://github.com/owcaelfis6626/spde under the MIT License. The repository includes the FD reference solver, vanilla and Spectral PINN implementations, FNO baseline, and analysis pipeline. Reproducing the main results requires approximately 8 hours on a single RTX 4090 GPU.

References

[1] R. Cont, P. Tankov, Financial Modelling with Jump Processes, Chapman & Hall/CRC, 2003.

[2] D. B. Madan, P. P. Carr, E. C. Chang, The variance gamma process and option pricing, European Finance Review 2 (1998) 79–105. doi:10.1023/A:1009703431535.

[3] M. Raissi, P. Perdikaris, G. E. Karniadakis, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, Journal of Computational Physics 378 (2019) 686–707. doi:10.1016/j.jcp.2018.10.045.

[4] L. Yang, X. Meng, G. E. Karniadakis, B-pinns: Bayesian physics-informed neural networks for forward and inverse pde problems with noisy data, Journal of Computational Physics 425 (2021) 109913. doi:10.1016/j.jcp.2020.109913.

[5] S. Mishra, R. Molinaro, Estimates on the generalization error of physics-informed neural networks for approximating pdes, IMA Journal of Numerical Analysis 42 (2) (2022) 981–1022. doi:10.1093/imanum/drab032.

[6] A. D. Jagtap, E. Kharazmi, G. E. Karniadakis, Conservative physics-informed neural networks on discrete domains for conservation laws, Journal of Computational Physics 422 (2020) 109762. doi:10.1016/j.jcp.2020.109762.

[7] E. Kharazmi, Z. Zhang, G. E. Karniadakis, Variational physics-informed neural networks for solving partial differential equations, arXiv preprint arXiv:1912.00873 (2019).

[8] S. Wang, X. Yu, P. Perdikaris, When and why pinns fail to train: A neural tangent kernel perspective, Journal of Computational Physics 449 (2022) 110768. doi:10.1016/j.jcp.2021.110768.

[9] S. Wang, S. Sankaran, P. Perdikaris, Respecting causality is all you need for training physics-informed neural networks, Computer Methods in Applied Mechanics and Engineering 421 (2024) 116813. doi:10.1016/j.cma.2024.116813.

[10] Z. Li, N. Kovachki, K. Azizzadenesheli, B. Liu, K. Bhattacharya, A. Stuart, A. Anandkumar, Fourier neural operator for parametric partial differential equations, in: International Conference on Learning Representations (ICLR), 2021.

[11] L. Lu, P. Jin, G. Pang, Z. Zhang, G. E. Karniadakis, Learning nonlinear operators via deeponet based on the universal approximation theorem of operators, Nature Machine Intelligence 3 (2021) 218–229. doi:10.1038/s42256-021-00302-5.

[12] Y. Du, N. Chalapathi, A. S. Krishnapriyan, Neural spectral methods: Self-supervised learning in the spectral domain, arXiv preprint arXiv:2312.05225 (2024).

[13] G. J. Lord, C. E. Powell, T. Shardlow, An Introduction to Computational Stochastic PDEs, Cambridge University Press, 2014. doi:10.1017/CBO9781139017329.

[14] G. Da Prato, J. Zabczyk, Stochastic Equations in Infinite Dimensions, 2nd Edition, Cambridge University Press, 2014. doi:10.1017/CBO9781107295513.

[15] M. Hairer, A theory of regularity structures, Inventiones mathematicae 198 (2) (2014) 269–504. doi:10.1007/s00222-014-0505-4.

[16] Z. Hu, Z. Zhang, G. E. Karniadakis, K. Kawaguchi, Score-fpinn: Score-based physics-informed neural networks for solving fokker–planck–lévy equations, arXiv preprint arXiv:2406.11676 (2024).

[17] D. E. Worrall, P. Lippe, J. Brandstetter, et al., Spectral shaping for neural pde surrogates, in: International Conference on Learning Representations (ICLR), 2025.

[18] M. Tancik, P. P. Srinivasan, B. Mildenhall, S. Fridovich-Keil, N. Raghavan, U. Singhal, R. Ramamoorthi, J. T. Barron, R. Ng, Fourier features let networks learn high frequency functions in low dimensional domains, in: Advances in Neural Information Processing Systems (NeurIPS), 2020.

[19] J. Sirignano, K. Spiliopoulos, Dgm: A deep learning algorithm for solving partial differential equations, Journal of Computational Physics 375 (2018) 1339–1364. doi:10.1016/j.jcp.2018.08.029.

[20] J. Kim, S. Lee, H. Park, Spectral-loss training for turbulence closure neural networks, in: NeurIPS Workshop on Machine Learning and the Physical Sciences, 2024.

[21] R. Durall, M. Keuper, J. Keuper, Watch your up-convolution: Cnn-based generative deep neural networks are failing to reproduce spectral distributions, in: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2020, pp. 7890–7899. doi:10.1109/CVPR42600.2020.00791.

[22] J. Engel, L. Hantrakul, C. Gu, A. Roberts, Ddsp: Differentiable digital signal processing, in: International Conference on Learning Representations (ICLR), 2020.

[23] D. W. Peaceman, H. H. Rachford, The numerical solution of parabolic and elliptic differential equations, Journal of the Society for Industrial and Applied Mathematics 3 (1) (1955) 28–41. doi:10.1137/0103003.

[24] G. B. Folland, Introduction to Partial Differential Equations, 2nd Edition, Princeton University Press, 1995.

[25] A. Pinkus, Approximation theory of the mlp model in neural networks, Acta Numerica 8 (1999) 143–195. doi:10.1017/S0962492900002919.

Appendix A. Detailed Proofs

Appendix B. Hyperparameter Tables

Appendix C. Additional Figures

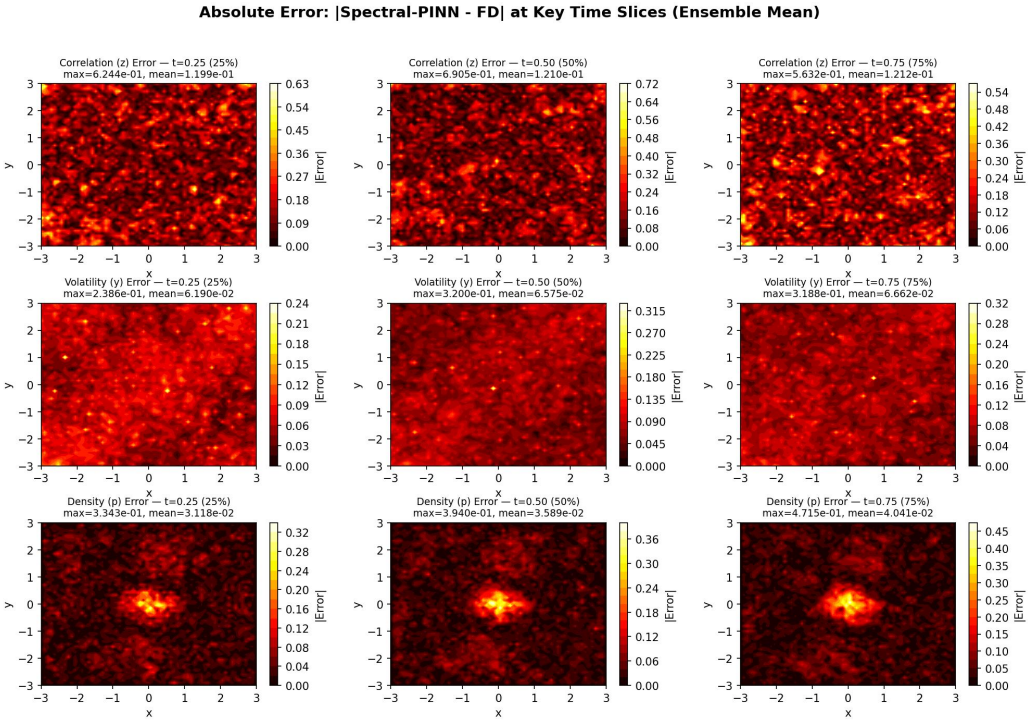

Figure C.4: Pointwise error heatmaps for the Spectral PINN vs. FD reference on each channel of the hierarchical 2D SPDE.

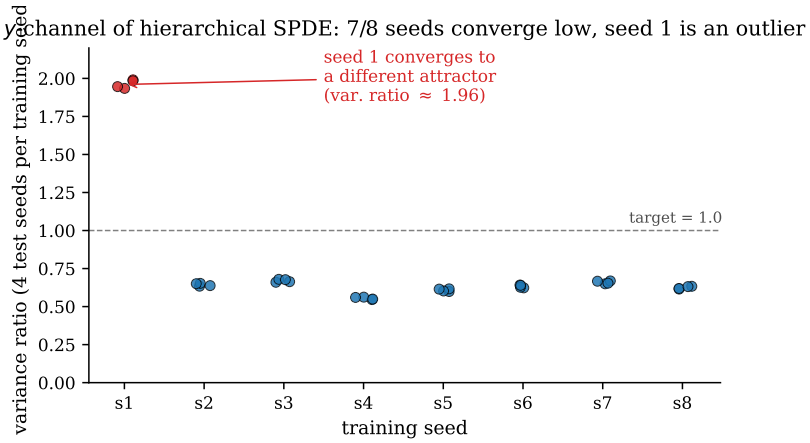

Figure C.5: Per-seed variance ratio on the y-channel of the hierarchical 2D SPDE. Seven of eight seeds converge to ≈ 0.63; seed 1 is a distinct attractor at ≈ 1.96.